

Heart_failure case study

May 15, 2024

0.0.1 Importing Libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

0.0.2 Loading Data

```
[4]: df=pd.read_csv('heart_failure_clinical_records.csv')
```

0.0.3 Exploring dataset

```
[6]: df.shape
```

```
[6]: (5000, 13)
```

```
[7]: df.head()
```

```
[7]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	\
0	55.0	0	748	0	45	
1	65.0	0	56	0	25	
2	45.0	0	582	1	38	
3	60.0	1	754	1	40	
4	95.0	1	582	0	30	

	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex	\
0	0	263358.03	1.3	137	1	
1	0	305000.00	5.0	130	1	
2	0	319000.00	0.9	140	0	
3	1	328000.00	1.2	126	1	
4	0	461000.00	2.0	132	1	

	smoking	time	DEATH_EVENT
0	1	88	0
1	0	207	0
2	0	244	0

3	0	90	0
4	0	50	1

```
[8]: df.tail()
```

```
[8]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	\
4995	45.0	0	582	1	55	
4996	60.0	1	582	0	30	
4997	95.0	1	112	0	40	
4998	65.0	1	160	1	20	
4999	40.0	0	244	0	45	

	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex	\
4995	0	543000.0	1.0	132	0	
4996	1	127000.0	0.9	145	0	
4997	1	196000.0	1.0	138	0	
4998	0	327000.0	2.7	116	0	
4999	1	275000.0	0.9	140	0	

	smoking	time	DEATH_EVENT
4995	0	250	0
4996	0	95	0
4997	0	24	1
4998	0	8	1
4999	0	174	0

```
[9]: df.columns
```

```
[9]: Index(['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes',
        'ejection_fraction', 'high_blood_pressure', 'platelets',
        'serum_creatinine', 'serum_sodium', 'sex', 'smoking', 'time',
        'DEATH_EVENT'],
        dtype='object')
```

```
[10]: df.info
```

```
[10]: <bound method DataFrame.info of
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	\
0	55.0	0	748	0	45	
1	65.0	0	56	0	25	
2	45.0	0	582	1	38	
3	60.0	1	754	1	40	
4	95.0	1	582	0	30	
...	
4995	45.0	0	582	1	55	
4996	60.0	1	582	0	30	
4997	95.0	1	112	0	40	

4998	65.0	1		160	1	20
4999	40.0	0		244	0	45

	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex	\
0	0	263358.03	1.3	137	1	
1	0	305000.00	5.0	130	1	
2	0	319000.00	0.9	140	0	
3	1	328000.00	1.2	126	1	
4	0	461000.00	2.0	132	1	
...	
4995	0	543000.00	1.0	132	0	
4996	1	127000.00	0.9	145	0	
4997	1	196000.00	1.0	138	0	
4998	0	327000.00	2.7	116	0	
4999	1	275000.00	0.9	140	0	

	smoking	time	DEATH_EVENT
0	1	88	0
1	0	207	0
2	0	244	0
3	0	90	0
4	0	50	1
...
4995	0	250	0
4996	0	95	0
4997	0	24	1
4998	0	8	1
4999	0	174	0

[5000 rows x 13 columns]>

0.0.4 Duplicated Values

```
[12]: df.duplicated().sum()
```

```
[12]: 3680
```

0.0.5 Missing Values

```
[14]: df.isna().sum()
```

```
[14]: age                0
anaemia                0
creatinine_phosphokinase  0
diabetes              0
ejection_fraction     0
```

```

high_blood_pressure      0
platelets                0
serum_creatinine         0
serum_sodium             0
sex                     0
smoking                  0
time                    0
DEATH_EVENT              0
dtype: int64

```

0.0.6 Data Distribution

```
[16]: df.dtypes
```

```

[16]: age                float64
      anaemia            int64
      creatinine_phosphokinase  int64
      diabetes           int64
      ejection_fraction  int64
      high_blood_pressure  int64
      platelets          float64
      serum_creatinine    float64
      serum_sodium        int64
      sex                int64
      smoking            int64
      time              int64
      DEATH_EVENT        int64
      dtype: object

```

0.0.7 Visualizing distribution

```

[18]: plt.figure(figsize=(15,10))
      for i, column in enumerate(df):
          plt.subplot(4,4,i+1)
          sns.histplot(df[column],bins=30,kde=True)
          plt.title(f'Distribution of {column}')
          plt.xlabel(column)
          plt.ylabel('Frequency')
          plt.xticks(rotation=45)
      plt.tight_layout()
      plt.show()

```

```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):

```

```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):

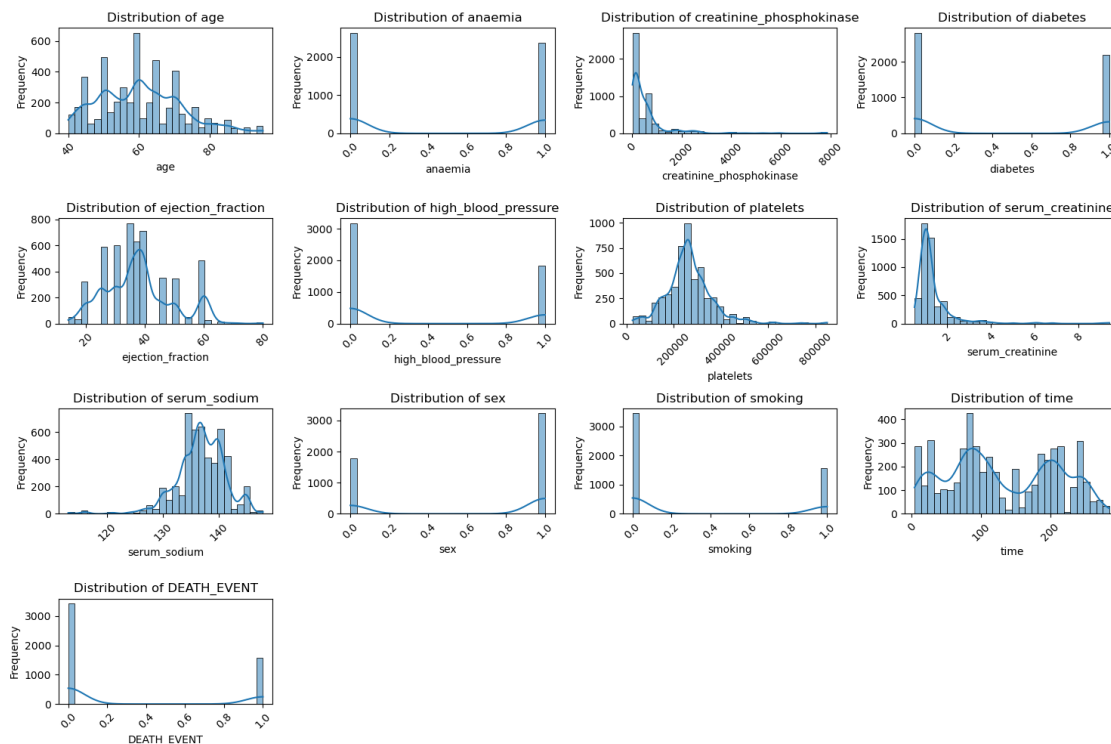
```

before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):  
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-  
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is  
deprecated and will be removed in a future version. Convert inf values to NaN  
before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-  
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is  
deprecated and will be removed in a future version. Convert inf values to NaN  
before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```



0.0.8 Visualizing skewness

```
[20]: from scipy import stats
```

```
[21]: plt.figure(figsize=(20, 15))  
for i, column in enumerate(df.columns):  
    plt.subplot(4, 4, i+1)  
    sns.histplot(df[column], bins=30, kde=True)  
    plt.title(f'Distribution of {column}')  
    plt.xlabel(column)
```

```

plt.ylabel('Frequency')

skewness = stats.skew(df[column])
if skewness < -1 or skewness > 1:
    plt.text(0.5, -0.2, f'Skewed ({skewness:.2f})',
    ↪horizontalalignment='center', verticalalignment='center', transform=plt.
    ↪gca().transAxes, fontsize=12)
else:
    plt.text(0.5, -0.2, 'Not Skewed', horizontalalignment='center',
    ↪verticalalignment='center', transform=plt.gca().transAxes, fontsize=12)
plt.tight_layout()
plt.show()

```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

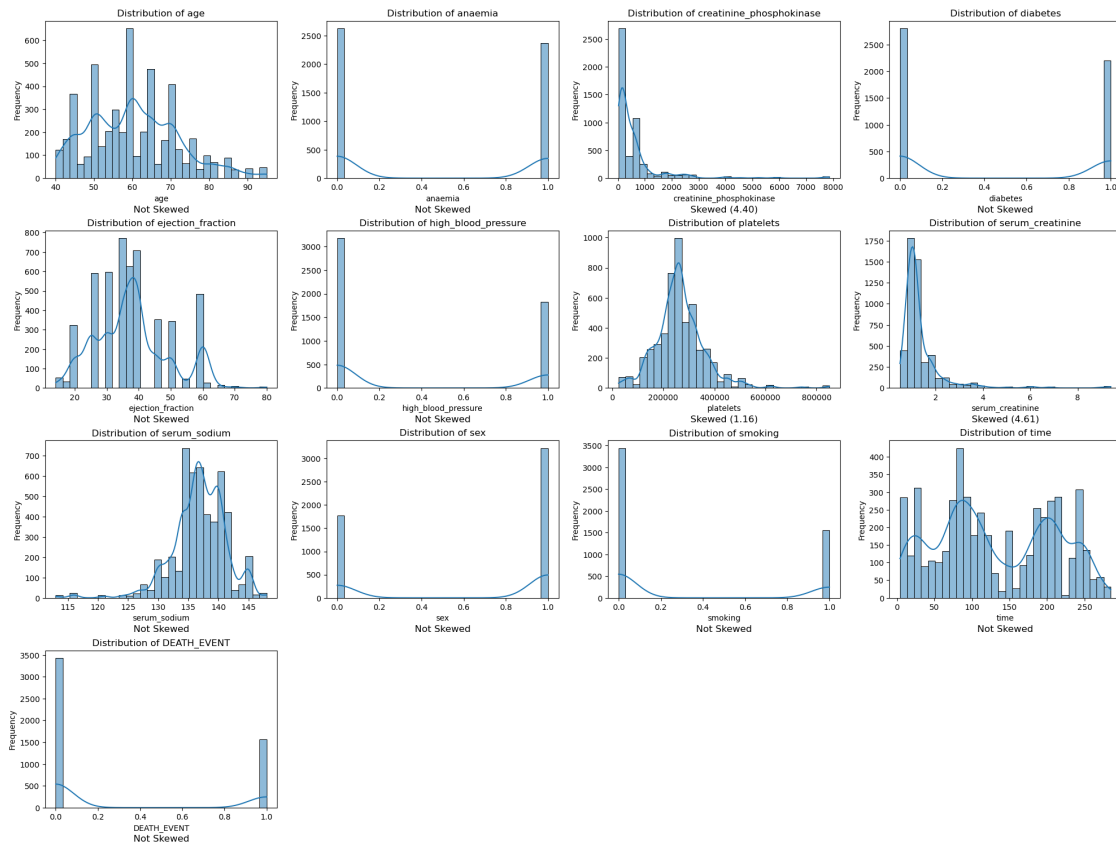
```
with pd.option_context('mode.use_inf_as_na', True):
```

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```

    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):

```

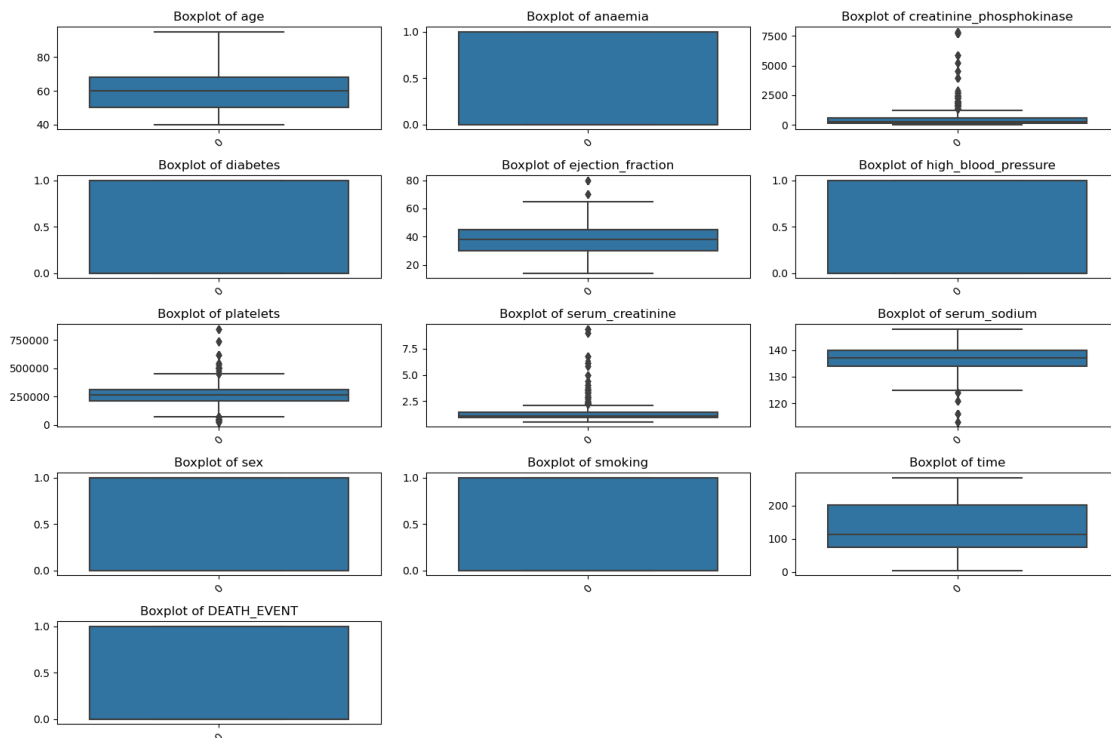
```
[22]: for col in df:
        skewness = df[col].skew()
        print(f"Skewness of '{col}': {skewness}")
```

```
Skewness of 'age': 0.45081168033242247
Skewness of 'anaemia': 0.10256525432936629
Skewness of 'creatinine_phosphokinase': 4.403444256118772
Skewness of 'diabetes': 0.2442735058883124
Skewness of 'ejection_fraction': 0.4944047809955734
Skewness of 'high_blood_pressure': 0.5618939376858764
Skewness of 'platelets': 1.157117070165634
Skewness of 'serum_creatinine': 4.615903417419207
Skewness of 'serum_sodium': -0.997974578554996
Skewness of 'sex': -0.6089660368778839
Skewness of 'smoking': 0.8128015889407144
Skewness of 'time': 0.11498892481750181
Skewness of 'DEATH_EVENT': 0.8037658645232266
```

0.1 Outliers

0.1.1 Checking for outliers

```
[25]: plt.figure(figsize=(15,10))
for i, column in enumerate(df):
    plt.subplot(5,3,i+1)
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Columns with outliers 'creatinine_phosphokinase', 'ejection_fraction', 'platelets', 'serum_creatinine', 'serum_sodium'

0.1.2 Percentage of outliers

```
[29]: def robust_zscore_outlier_removal(df, columns):
    for column in columns:
        median = df[column].median()
        mad = np.median(np.abs(df[column] - median))
        modified_z_scores = 0.6745 * (df[column] - median) / mad
        df = df[(modified_z_scores > -3.5) & (modified_z_scores < 3.5)]
```

```

return df

outlier_columns = ['creatinine_phosphokinase', 'ejection_fraction',
↳ 'platelets', 'serum_creatinine', 'serum_sodium']
df1 = robust_zscore_outlier_removal(df, outlier_columns)
def calculate_outlier_percentages(df, columns):
    outlier_percentages = {}

    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]

        outlier_percentage = len(outliers) / len(df) * 100
        outlier_percentages[column] = outlier_percentage
    return outlier_percentages
outlier_percentages_df1 = calculate_outlier_percentages(df1, outlier_columns)
print("Percentage of outliers in each column of df1 after robust Z-score_
↳ cleaning:")
for column, percentage in outlier_percentages_df1.items():
    print(f"{column}: {percentage:.2f}%")

```

Percentage of outliers in each column of df1 after robust Z-score cleaning:

```

creatinine_phosphokinase: 0.00%
ejection_fraction: 0.00%
platelets: 5.10%
serum_creatinine: 10.60%
serum_sodium: 2.13%

```

0.1.3 Handling outliers using quartile method

```

[31]: def robust_zscore_outlier_removal(df, columns, thresholds):
    for column in columns:
        median = df[column].median()
        mad = np.median(np.abs(df[column] - median))
        modified_z_scores = 0.6745 * (df[column] - median) / mad
        threshold = thresholds[column]
        df = df[(modified_z_scores > -threshold) & (modified_z_scores <_
↳ threshold)]

    return df
outlier_columns = ['creatinine_phosphokinase', 'ejection_fraction',
↳ 'platelets', 'serum_creatinine', 'serum_sodium']

```

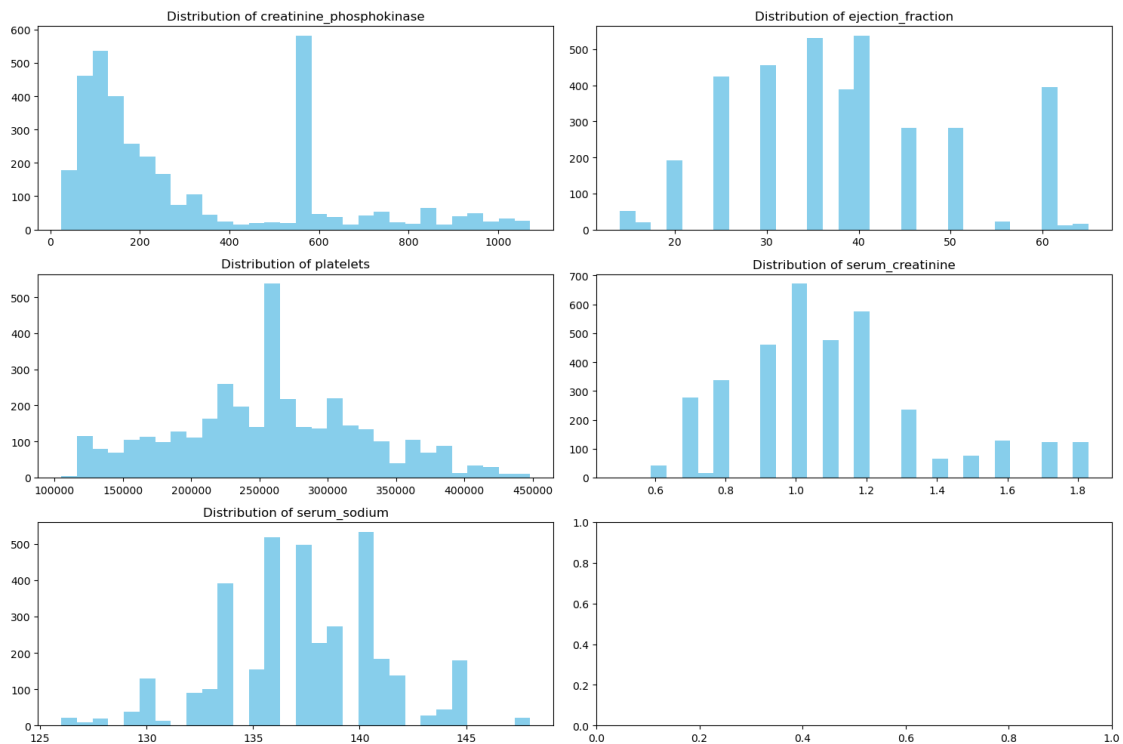
```

thresholds = {
    'creatinine_phosphokinase': 3.5, # No outliers, keep standard
    'ejection_fraction': 3.5,      # No outliers, keep standard
    'platelets': 2.5,               # More stringent due to persistent
    ↪outliers
    'serum_creatinine': 2.5,        # More stringent due to persistent
    ↪outliers
    'serum_sodium': 2.5             # More stringent due to persistent
    ↪outliers
}

df1 = robust_zscore_outlier_removal(df, outlier_columns, thresholds)
fig, axs = plt.subplots(3, 2, figsize=(15, 10))
axs = axs.flatten()
for i, column in enumerate(outlier_columns):
    axs[i].hist(df1[column], bins=30, color='skyblue')
    axs[i].set_title(f'Distribution of {column}')
plt.tight_layout()
plt.show()
outlier_percentages_df1 = calculate_outlier_percentages(df1, outlier_columns)

print("Percentage of outliers in each column of df1 after adjusted robust
    ↪Z-score cleaning:")
for column, percentage in outlier_percentages_df1.items():
    print(f"{column}: {percentage:.2f}%")

```



Percentage of outliers in each column of df1 after adjusted robust Z-score cleaning:

creatinine_phosphokinase: 0.00%

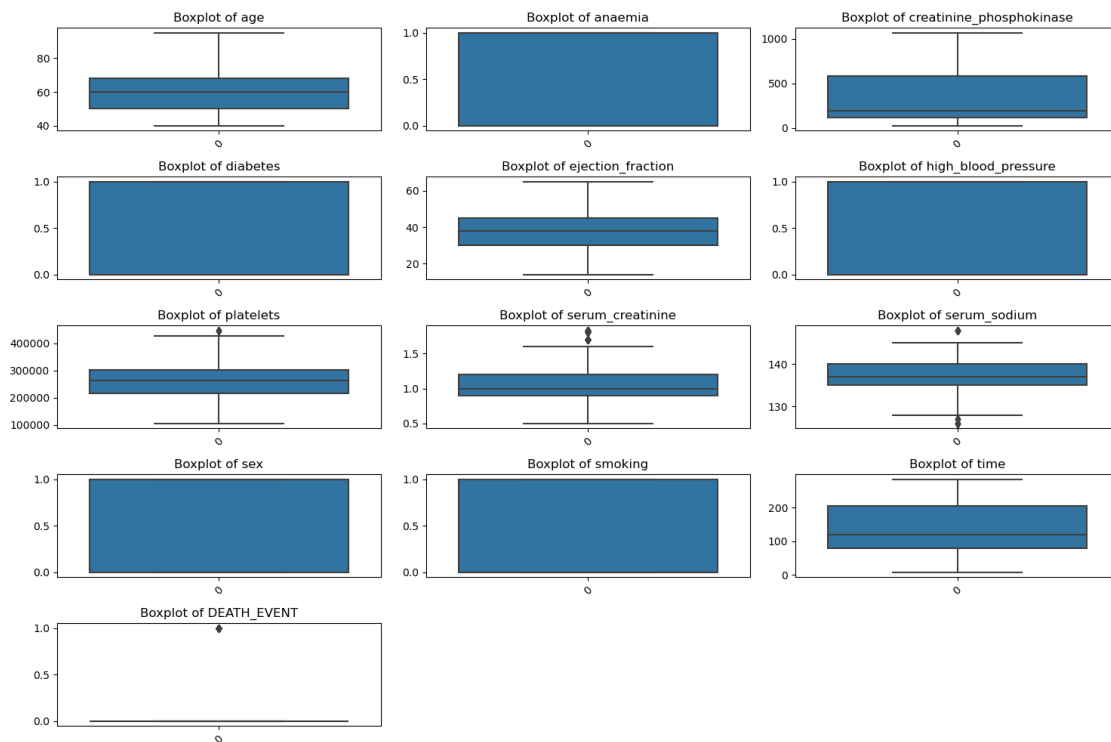
ejection_fraction: 0.00%

platelets: 0.30%

serum_creatinine: 6.79%

serum_sodium: 1.44%

```
[40]: plt.figure(figsize=(15,10))
for i, column in enumerate(df1):
    plt.subplot(5,3,i+1)
    sns.boxplot(df1[column])
    plt.title(f'Boxplot of {column}')
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



0.1.4 Identifying the best model

```
[50]: from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB, GaussianNB
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix

features = ['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes', \
    'ejection_fraction', 'high_blood_pressure', \
    'platelets', 'serum_creatinine', 'serum_sodium', 'sex', 'smoking', \
    'time']
target = 'DEATH_EVENT'
X = df1[features]
y = df1[target]

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state=42)

# List of classifiers
classifiers = {
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(),
    "KNN": KNeighborsClassifier(),
    "Logistic Regression": LogisticRegression(),
    "Multinomial Naive Bayes": MultinomialNB(),
    "Gaussian Naive Bayes": GaussianNB()
}

# Training and evaluating each classifier
results = {}
for name, clf in classifiers.items():
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    results[name] = {
        "Accuracy": accuracy,
        "Misclassification Rate": 1 - accuracy,
        "Classification Report": classification_report(y_test, y_pred),
        "Confusion Matrix": confusion_matrix(y_test, y_pred)
    }
```

```

for name, result in results.items():
    print(f"{name}:")
    print(f"Accuracy: {result['Accuracy']}")
    print(f"Misclassification Rate: {result['Misclassification Rate']}")
    print("Classification Report:")
    print(result["Classification Report"])
    print("Confusion Matrix:")
    print(result["Confusion Matrix"])
    print()

```

Decision Tree:

Accuracy: 0.9847645429362881

Misclassification Rate: 0.015235457063711877

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	531
1	0.98	0.96	0.97	191
accuracy			0.98	722
macro avg	0.98	0.98	0.98	722
weighted avg	0.98	0.98	0.98	722

Confusion Matrix:

```

[[528  3]
 [ 8 183]]

```

Random Forest:

Accuracy: 0.9916897506925207

Misclassification Rate: 0.008310249307479256

Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	531
1	0.99	0.97	0.98	191
accuracy			0.99	722
macro avg	0.99	0.99	0.99	722
weighted avg	0.99	0.99	0.99	722

Confusion Matrix:

```

[[530  1]
 [ 5 186]]

```

SVM:

Accuracy: 0.7354570637119113

Misclassification Rate: 0.26454293628808867

Classification Report:

	precision	recall	f1-score	support
0	0.74	1.00	0.85	531
1	0.00	0.00	0.00	191
accuracy			0.74	722
macro avg	0.37	0.50	0.42	722
weighted avg	0.54	0.74	0.62	722

Confusion Matrix:

```
[[531  0]
 [191  0]]
```

KNN:

Accuracy: 0.9626038781163435

Misclassification Rate: 0.037396121883656486

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.98	0.97	531
1	0.95	0.91	0.93	191
accuracy			0.96	722
macro avg	0.96	0.95	0.95	722
weighted avg	0.96	0.96	0.96	722

Confusion Matrix:

```
[[521 10]
 [ 17 174]]
```

Logistic Regression:

Accuracy: 0.8601108033240997

Misclassification Rate: 0.13988919667590027

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.95	0.91	531
1	0.83	0.60	0.69	191
accuracy			0.86	722
macro avg	0.85	0.78	0.80	722
weighted avg	0.86	0.86	0.85	722

Confusion Matrix:

```
[[507 24]
 [ 77 114]]
```


Multinomial Naive Bayes:

Accuracy: 0.724376731301939

Misclassification Rate: 0.275623268698061

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.75	0.80	531
1	0.48	0.66	0.56	191
accuracy			0.72	722
macro avg	0.67	0.70	0.68	722
weighted avg	0.76	0.72	0.74	722

Confusion Matrix:

```
[[397 134]
```

```
[ 65 126]]
```

Gaussian Naive Bayes:

Accuracy: 0.8795013850415513

Misclassification Rate: 0.12049861495844871

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.97	0.92	531
1	0.89	0.62	0.73	191
accuracy			0.88	722
macro avg	0.88	0.80	0.83	722
weighted avg	0.88	0.88	0.87	722

Confusion Matrix:

```
[[516 15]
```

```
[ 72 119]]
```

```
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

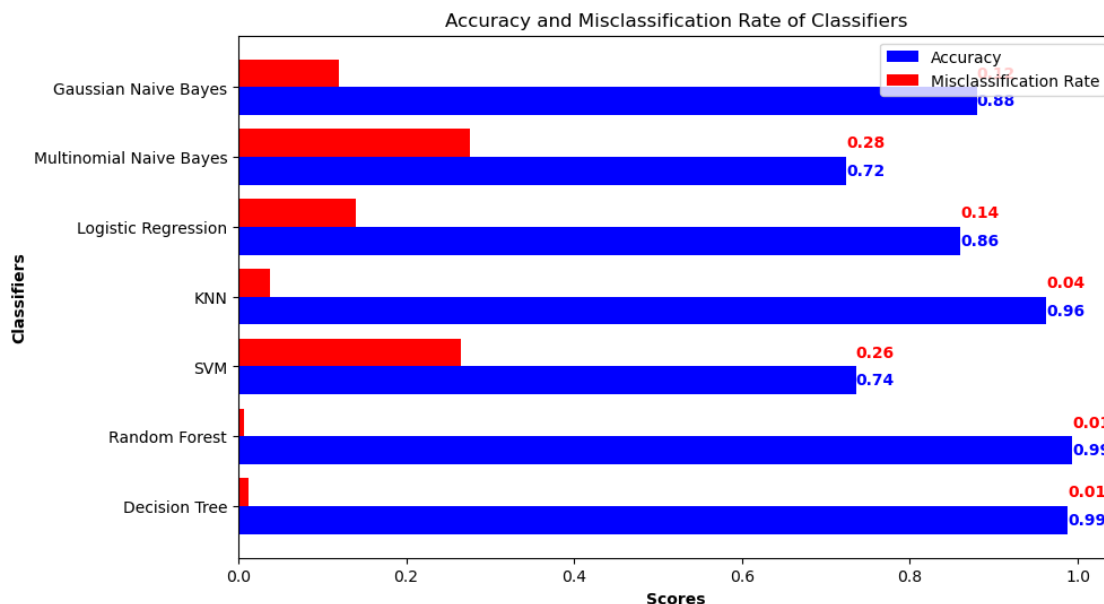
```
[54]: fig, ax = plt.subplots(figsize=(10, 6))
      y_pos = np.arange(len(classifiers_names))

      # Plotting bars
      bar1 = ax.barh(y_pos, accuracy_values, color='b', height=0.4, label='Accuracy')
      bar2 = ax.barh(y_pos + 0.4, misclassification_rates, color='r', height=0.4,
                    label='Misclassification Rate')

      # Adding labels
      ax.set_xlabel('Scores', fontweight='bold')
      ax.set_ylabel('Classifiers', fontweight='bold')
      ax.set_yticks(y_pos + 0.2)
      ax.set_yticklabels(classifiers_names)
      ax.legend()

      # Adding labels to bars
      for bar, acc, mis in zip(bar1, accuracy_values, misclassification_rates):
          ax.text(bar.get_width(), bar.get_y() + bar.get_height() / 2, f'{acc:.2f}',
                  va='center', ha='left', fontsize=10, fontweight='bold', color='blue')
          ax.text(bar.get_width(), bar.get_y() + bar.get_height() / 2 + 0.4, f'{mis:.2f}',
                  va='center', ha='left', fontsize=10, fontweight='bold', color='red')

      # Displaying the plot
      plt.title('Accuracy and Misclassification Rate of Classifiers')
      plt.show()
```



0.1.5 Using Random Forest

```
[65]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score

      # Features and target variable
      features = ['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes',
                  ↪ 'ejection_fraction', 'high_blood_pressure',
                  ↪ 'platelets', 'serum_creatinine', 'serum_sodium', 'sex', 'smoking',
                  ↪ 'time']
      target = 'DEATH_EVENT'
      X = df1[features]
      y = df1[target]

      # Splitting the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                  ↪ random_state=42)
      rf_classifier = RandomForestClassifier()
      rf_classifier.fit(X_train, y_train)

      # Predictions on the test set
      y_pred_rf = rf_classifier.predict(X_test)

      # Accuracy
      accuracy_rf = accuracy_score(y_test, y_pred_rf)

      # Misclassification Rate
      misclassification_rate_rf = 1 - accuracy_rf

      print("Random Forest:")
      print(f"Accuracy: {accuracy_rf}")
      print(f"Misclassification Rate: {misclassification_rate_rf}")
```

Random Forest:

Accuracy: 0.9916897506925207

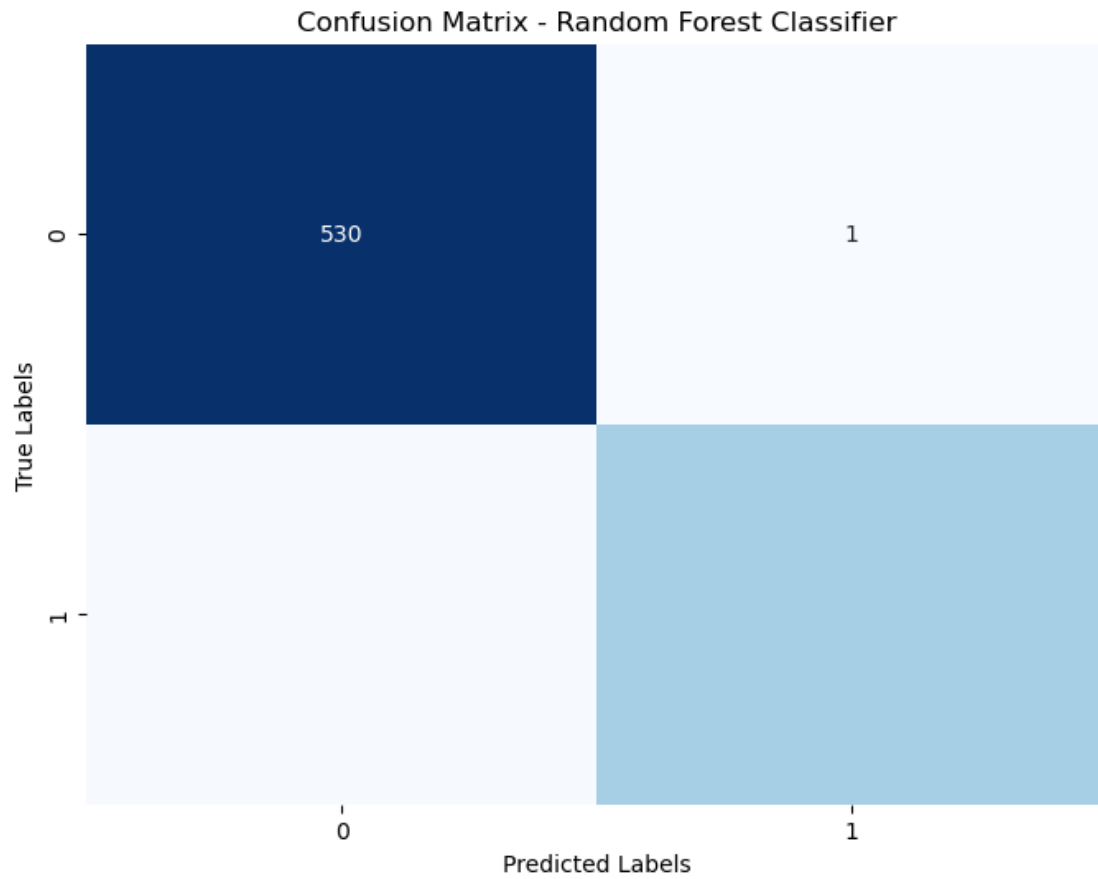
Misclassification Rate: 0.008310249307479256

0.1.6 Confusion Matrix

```
[68]: from sklearn.metrics import confusion_matrix
      conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

      plt.figure(figsize=(8, 6))
      sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Blues', cbar=False)
```

```
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Random Forest Classifier')
plt.show()
```



[]: